

advance_cross_prediction

August 23, 2021

0.1 Import Libraries

```
[1]: import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
# %matplotlib inline
import seaborn as sns

# pd.set_option('display.max_columns',30)
# pd.set_option('display.max_rows',500)
```

0.2 Load the Dataset

```
[2]: train_data=pd.read_csv('kwope/Train/train.csv')
test_data = pd.read_csv('kwope/Test/Test.csv')
```

```
[3]: train_data.head()
```

```
[3]:
```

	id	Gender	Age	Driving_License	Region_Code	Previously_Insured	\
0	274711	Male	22	1	40.0	0	
1	216540	Female	26	1	8.0	1	
2	78984	Female	32	1	28.0	0	
3	218369	Male	32	1	37.0	0	
4	307216	Female	24	1	28.0	0	

	Vehicle_Age	Vehicle_Damage	Policy_Sales_Channel	Vintage	Annual_Premium
0	< 1 Year	Yes	152.0	203	25455.0
1	< 1 Year	No	26.0	45	38393.0
2	< 1 Year	Yes	122.0	35	32118.0
3	< 1 Year	Yes	152.0	215	37711.0
4	< 1 Year	Yes	156.0	222	48706.0

```
[4]: test_data.head()
```

```
[4]:
```

	id	Gender	Age	Driving_License	Region_Code	Previously_Insured	\
0	16222	Male	54	1	28.0	0	

1	342481	Female	26	1	3.0	0
2	20916	Male	25	1	28.0	0
3	38704	Male	29	1	24.0	0
4	363809	Female	28	1	28.0	0

	Vehicle_Age	Vehicle_Damage	Policy_Sales_Channel	Vintage
0	> 2 Years	Yes	26.0	53
1	1-2 Year	Yes	156.0	280
2	< 1 Year	Yes	124.0	255
3	1-2 Year	Yes	157.0	235
4	< 1 Year	Yes	26.0	243

Observation * 'Annual_Premium' only is omitted in test_data

Description of each of the column headers in the data frame id: randomly generated numbers for customer

Gender: Gender of customer

Age: age of customer

Driving_licence: if customer has driving licence or not (yes=1, no=0)

Region_Code: region code of customer

Previously_Insured: if customer was previously insured or not (yes=1, no=0)

Vehicle_Age: age of customer's vehicle (< 1 year, 1-2 year, or >2 years)

Vehicle_Damage: if customer's vehicle has been damaged or not (yes=1, no=0)

Policy_sales_channel: No. of sales channel

vintage: vintage

Annual_Premium: Annual premium bought by customer

0.3 Exploratory Data Analysis (EDA)

```
[5]: #Get the shape of the datasets
print('train size', train_data.shape)
print('train size', test_data.shape)
```

train size (304887, 11)

train size (76222, 10)

```
[6]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 304887 entries, 0 to 304886
Data columns (total 11 columns):
#   Column                Non-Null Count  Dtype
---  -
#   Column                Non-Null Count  Dtype
```

```

0   id                304887 non-null int64
1   Gender             304887 non-null object
2   Age               304887 non-null int64
3   Driving_License    304887 non-null int64
4   Region_Code        304887 non-null float64
5   Previously_Insured 304887 non-null int64
6   Vehicle_Age        304887 non-null object
7   Vehicle_Damage     304887 non-null object
8   Policy_Sales_Channel 304887 non-null float64
9   Vintage            304887 non-null int64
10  Annual_Premium     304887 non-null float64

```

dtypes: float64(3), int64(5), object(3)

memory usage: 25.6+ MB

```
[7]: #get a statistical summary of the dataset
train_data.describe()
```

```
[7]:
```

	id	Age	Driving_License	Region_Code	\
count	304887.000000	304887.000000	304887.000000	304887.000000	
mean	190738.657112	38.826897	0.997855	26.396239	
std	110004.367239	15.515299	0.046265	13.228749	
min	1.000000	20.000000	0.000000	0.000000	
25%	95504.500000	25.000000	1.000000	15.000000	
50%	190886.000000	36.000000	1.000000	28.000000	
75%	285863.500000	49.000000	1.000000	35.000000	
max	381109.000000	85.000000	1.000000	52.000000	

	Previously_Insured	Policy_Sales_Channel	Vintage	Annual_Premium
count	304887.000000	304887.000000	304887.000000	304887.000000
mean	0.458127	112.053859	154.392214	30591.308311
std	0.498244	54.189288	83.670312	17239.285689
min	0.000000	1.000000	10.000000	2630.000000
25%	0.000000	29.000000	82.000000	24403.000000
50%	0.000000	134.000000	154.000000	31697.000000
75%	1.000000	152.000000	227.000000	39443.000000
max	1.000000	163.000000	299.000000	540165.000000

```
[8]: train_data.nunique()
```

```
[8]: id                304887
Gender                2
Age                  66
Driving_License      2
Region_Code          53
Previously_Insured    2
Vehicle_Age           3
Vehicle_Damage        2
Policy_Sales_Channel 154
```

```
Vintage                290
Annual_Premium        46479
dtype: int64
```

```
[9]: # Finding count of null values in each of the columns in data
train_data.isnull().sum()
```

```
[9]: id                0
Gender                0
Age                  0
Driving_License       0
Region_Code           0
Previously_Insured     0
Vehicle_Age           0
Vehicle_Damage         0
Policy_Sales_Channel   0
Vintage              0
Annual_Premium         0
dtype: int64
```

```
[10]: test_data.isnull().sum()
```

```
[10]: id                0
Gender                0
Age                  0
Driving_License       0
Region_Code           0
Previously_Insured     0
Vehicle_Age           0
Vehicle_Damage         0
Policy_Sales_Channel   0
Vintage              0
dtype: int64
```


Observation * No Null Values in both train and test dataset * 'Gender', 'Vehicle_Age', and 'Vehicle_Damage' columns of 'object' dtype . so we need check the unique values of those columns, if any string values is present.

Observation

'Gender', 'Vehicle_Age', 'Vehicle_Damage' columns have few string values. We need to replace the string with corresponding int values.

```
[11]: # sns.pairplot(train_data[1000:-1])
sns.displot(train_data['Annual_Premium'])
# dt = sns.load_dataset('kowope/Train/train.csv')
# dt.head()
```

```
[11]: <seaborn.axisgrid.FacetGrid at 0x1bf7a13d6a0>
```




```
[12]: plt.figure(figsize=(16,16))
for i, col in enumerate(["id", "Age", "
    ↳ "Region_Code", "Annual_Premium", "Policy_Sales_Channel", "Vintage"]):
    plt.subplot(4,4,i+1)
    sns.boxplot(train_data[col])
    plt.tight_layout()
```

C:\Users\hp\Anaconda3\envs\tf\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.

warnings.warn(
C:\Users\hp\Anaconda3\envs\tf\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.

```
warnings.warn(
C:\Users\hp\Anaconda3\envs\tf\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(
C:\Users\hp\Anaconda3\envs\tf\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(
C:\Users\hp\Anaconda3\envs\tf\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(
C:\Users\hp\Anaconda3\envs\tf\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(
```



```
[13]: # check if id values are unique or not
for i in train_data['id']:
```

```
if i==test_data['id'][67]:
    print('yes')
```

```
[14]: print("Unique Values in 'Gender' column are ",train_data['Gender'].unique())
      print("Unique Values in 'Vehicle_Age' column are ",train_data['Vehicle_Age'].
      ↪unique())
      print("Unique Values in 'Vehicle_Damage' column are_
      ↪",train_data['Vehicle_Damage'].unique())
```

```
Unique Values in 'Gender' column are  ['Male' 'Female']
Unique Values in 'Vehicle_Age' column are  ['< 1 Year' '1-2 Year' '> 2 Years']
Unique Values in 'Vehicle_Damage' column are  ['Yes' 'No']
```

```
[15]: # Checking frequency of each values in categorical columns_
      ↪'job_type','marital_status','education_level','communication'
      # & 'outcome' columns
      train_data.Gender.value_counts()
```

```
[15]: Male      164780
      Female    140107
      Name: Gender, dtype: int64
```

```
[16]: train_data.Vehicle_Age.value_counts()
```

```
[16]: 1-2 Year      160195
      < 1 Year     131852
      > 2 Years    12840
      Name: Vehicle_Age, dtype: int64
```

```
[17]: train_data.Vehicle_Damage.value_counts()
```

```
[17]: Yes       153959
      No        150928
      Name: Vehicle_Damage, dtype: int64
```

```
[18]: train_data.Driving_License.value_counts()
```

```
[18]: 1      304233
      0       654
      Name: Driving_License, dtype: int64
```

0.4 Data Cleaning and Data Preprocessing

0.4.1 Converting Non Numeric Data Columns to Numeric Columns

As the model built for prediction understands numerical data columns better, we need to convert categorical data columns to numerical data columns by using dummy columns and one hot encoding methods

```
[19]: mapping_dict = {
      "Gender": {
          "Male": 0,
          "Female": 1
      },
      "Vehicle_Damage": {
          "No": 0,
          "Yes": 1
      },
      # "Vehicle_Age": {
      #     "< 1 Year": 0,
      #     "1-2 Year": 1,
      #     "> 2 Years": 2,
      # }
    }
```

```
[20]: filtered_train_data = train_data.replace(mapping_dict)
      filtered_train_data[['Gender', 'Vehicle_Damage']]
      filtered_train_data = pd.get_dummies(filtered_train_data,
      ↪columns=['Vehicle_Age'])
```

```
[21]: filtered_test_data = test_data.replace(mapping_dict)
      filtered_test_data[['Gender', 'Vehicle_Damage']]
      filtered_test_data = pd.get_dummies(filtered_test_data, columns=['Vehicle_Age'])
```

```
[22]: filtered_train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 304887 entries, 0 to 304886
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     304887 non-null  int64
1   Gender                                304887 non-null  int64
2   Age                                   304887 non-null  int64
3   Driving_License                       304887 non-null  int64
4   Region_Code                           304887 non-null  float64
5   Previously_Insured                    304887 non-null  int64
6   Vehicle_Damage                        304887 non-null  int64
7   Policy_Sales_Channel                  304887 non-null  float64
8   Vintage                               304887 non-null  int64
9   Annual_Premium                        304887 non-null  float64
10  Vehicle_Age_1-2 Year                  304887 non-null  uint8
11  Vehicle_Age_< 1 Year                  304887 non-null  uint8
12  Vehicle_Age_> 2 Years                 304887 non-null  uint8
dtypes: float64(3), int64(7), uint8(3)
memory usage: 24.1 MB
```



```
[23]: filtered_train_data = filtered_train_data.rename(columns={
    "Vehicle_Age_1-2 Year": "Vehicle_Age_1_2_Year",
    "Vehicle_Age_< 1 Year": "Vehicle_Age_lt_1_Year",
    "Vehicle_Age_> 2 Years": "Vehicle_Age_gt_2_Years"
})
filtered_train_data["Vehicle_Age_lt_1_Year"] = filtered_train_data["Vehicle_Age_lt_1_Year"].
    ↪ astype('int')
filtered_train_data["Vehicle_Age_1_2_Year"] = filtered_train_data["Vehicle_Age_1_2_Year"].
    ↪ astype('int')
filtered_train_data["Vehicle_Age_gt_2_Years"] = filtered_train_data["Vehicle_Age_gt_2_Years"].
    ↪ astype('int')
filtered_train_data["Vehicle_Damage"] = filtered_train_data["Vehicle_Damage"].
    ↪ astype('int')
```

```
[24]: filtered_test_data = filtered_test_data.rename(columns={
    "Vehicle_Age_1-2 Year": "Vehicle_Age_1_2_Year",
    "Vehicle_Age_< 1 Year": "Vehicle_Age_lt_1_Year",
    "Vehicle_Age_> 2 Years": "Vehicle_Age_gt_2_Years"
})
filtered_test_data["Vehicle_Age_lt_1_Year"] = filtered_test_data["Vehicle_Age_lt_1_Year"].
    ↪ astype('int')
filtered_test_data["Vehicle_Age_1_2_Year"] = filtered_test_data["Vehicle_Age_1_2_Year"].
    ↪ astype('int')
filtered_test_data["Vehicle_Age_gt_2_Years"] = filtered_test_data["Vehicle_Age_gt_2_Years"].
    ↪ astype('int')
filtered_test_data["Vehicle_Damage"] = filtered_test_data["Vehicle_Damage"].
    ↪ astype('int')
```

0.4.2 removing outliers

```
[25]: filtered_train_data.loc[filtered_train_data.Annual_Premium > 400000,
    ↪ "Annual_Premium"] = 400000
```

```
[26]: filtered_train_data = filtered_train_data.drop(['id'], axis=1)
# filtered_train_data.dropna()
```

```
[27]: filtered_train_data.head()
```

```
[27]:
```

	Gender	Age	Driving_License	Region_Code	Previously_Insured	\
0	0	22	1	40.0	0	
1	1	26	1	8.0	1	
2	1	32	1	28.0	0	
3	0	32	1	37.0	0	
4	1	24	1	28.0	0	

	Vehicle_Damage	Policy_Sales_Channel	Vintage	Annual_Premium	\
0	1	152.0	203	25455.0	

1	0	26.0	45	38393.0
2	1	122.0	35	32118.0
3	1	152.0	215	37711.0
4	1	156.0	222	48706.0

	Vehicle_Age_1_2_Year	Vehicle_Age_lt_1_Year	Vehicle_Age_gt_2_Years
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0

```
[28]: filtered_train_data.describe()
```

```
[28]:
```

	Gender	Age	Driving_License	Region_Code \
count	304887.000000	304887.000000	304887.000000	304887.000000
mean	0.459537	38.826897	0.997855	26.396239
std	0.498361	15.515299	0.046265	13.228749
min	0.000000	20.000000	0.000000	0.000000
25%	0.000000	25.000000	1.000000	15.000000
50%	0.000000	36.000000	1.000000	28.000000
75%	1.000000	49.000000	1.000000	35.000000
max	1.000000	85.000000	1.000000	52.000000

	Previously_Insured	Vehicle_Damage	Policy_Sales_Channel \
count	304887.000000	304887.000000	304887.000000
mean	0.458127	0.504971	112.053859
std	0.498244	0.499976	54.189288
min	0.000000	0.000000	1.000000
25%	0.000000	0.000000	29.000000
50%	0.000000	1.000000	134.000000
75%	1.000000	1.000000	152.000000
max	1.000000	1.000000	163.000000

	Vintage	Annual_Premium	Vehicle_Age_1_2_Year \
count	304887.000000	304887.000000	304887.000000
mean	154.392214	30588.781565	0.525424
std	83.670312	17177.346735	0.499354
min	10.000000	2630.000000	0.000000
25%	82.000000	24403.000000	0.000000
50%	154.000000	31697.000000	1.000000
75%	227.000000	39443.000000	1.000000
max	299.000000	400000.000000	1.000000

	Vehicle_Age_lt_1_Year	Vehicle_Age_gt_2_Years
count	304887.000000	304887.000000
mean	0.432462	0.042114

std	0.495418	0.200849
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	0.000000
75%	1.000000	0.000000
max	1.000000	1.000000

0.5 Build Machine Learning Model & Evaluate it

```
[29]: from sklearn.ensemble import RandomForestRegressor
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.linear_model import Lasso, LinearRegression
from sklearn.model_selection import \
    train_test_split, GridSearchCV, RandomizedSearchCV
from sklearn.metrics import classification_report, confusion_matrix, \
    mean_squared_error, mean_absolute_error
import xgboost as xgb
```

```
[30]: #feature scaling
#normalize the "Age" and "Vintage" features with standardization
ss = StandardScaler()
num_feat = ['Age', 'Vintage']
filtered_train_data[num_feat] = ss.fit_transform(filtered_train_data[num_feat])
filtered_test_data[num_feat] = ss.fit_transform(filtered_test_data[num_feat])

#normalize "Annual_Premium" with min-max scaling
mm = MinMaxScaler()
filtered_train_data[['Annual_Premium']] = mm.
    fit_transform(filtered_train_data[['Annual_Premium']])
filtered_train_data
```

```
[30]:
```

	Gender	Age	Driving_License	Region_Code	Previously_Insured	\
0	0	-1.084538	1	40.0	0	
1	1	-0.826727	1	8.0	1	
2	1	-0.440011	1	28.0	0	
3	0	-0.440011	1	37.0	0	
4	1	-0.955632	1	28.0	0	
...	
304882	0	-0.891180	1	47.0	1	
304883	0	1.944736	1	28.0	0	
304884	0	-0.955632	1	28.0	0	
304885	1	0.075610	1	8.0	1	
304886	0	-0.955632	1	15.0	0	

	Vehicle_Damage	Policy_Sales_Channel	Vintage	Annual_Premium	\
0	1	152.0	0.580945	0.057440	
1	0	26.0	-1.307422	0.089999	

2	1	122.0	-1.426939	0.074208
3	1	152.0	0.724365	0.088283
4	1	156.0	0.808027	0.115952
...
304882	1	160.0	-1.152050	0.051287
304883	1	26.0	1.620742	0.072980
304884	0	152.0	-1.104243	0.142094
304885	0	26.0	1.286095	0.066210
304886	1	152.0	1.214385	0.065058

	Vehicle_Age_1_2_Year	Vehicle_Age_lt_1_Year	Vehicle_Age_gt_2_Years
0	0	1	0
1	0	1	0
2	0	1	0
3	0	1	0
4	0	1	0
...
304882	0	1	0
304883	1	0	0
304884	0	1	0
304885	1	0	0
304886	0	1	0

[304887 rows x 12 columns]

```
[31]: filtered_train_data.describe()
```

```
[31]:
```

	Gender	Age	Driving_License	Region_Code \
count	304887.000000	3.048870e+05	304887.000000	304887.000000
mean	0.459537	2.169590e-16	0.997855	26.396239
std	0.498361	1.000002e+00	0.046265	13.228749
min	0.000000	-1.213443e+00	0.000000	0.000000
25%	0.000000	-8.911797e-01	1.000000	15.000000
50%	0.000000	-1.822009e-01	1.000000	28.000000
75%	1.000000	6.556832e-01	1.000000	35.000000
max	1.000000	2.975977e+00	1.000000	52.000000

	Previously_Insured	Vehicle_Damage	Policy_Sales_Channel	Vintage \
count	304887.000000	304887.000000	304887.000000	3.048870e+05
mean	0.458127	0.504971	112.053859	-6.448526e-17
std	0.498244	0.499976	54.189288	1.000002e+00
min	0.000000	0.000000	1.000000	-1.725731e+00
25%	0.000000	0.000000	29.000000	-8.652093e-01
50%	0.000000	1.000000	134.000000	-4.687622e-03
75%	1.000000	1.000000	152.000000	8.677858e-01
max	1.000000	1.000000	163.000000	1.728307e+00

	Annual_Premium	Vehicle_Age_1_2_Year	Vehicle_Age_lt_1_Year \
count	304887.000000	304887.000000	304887.000000
mean	0.070360	0.525424	0.432462
std	0.043228	0.499354	0.495418
min	0.000000	0.000000	0.000000
25%	0.054793	0.000000	0.000000
50%	0.073148	1.000000	0.000000
75%	0.092642	1.000000	1.000000
max	1.000000	1.000000	1.000000

	Vehicle_Age_gt_2_Years
count	304887.000000
mean	0.042114
std	0.200849
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	1.000000

```
[32]: X = filtered_train_data.drop(['Annual_Premium'], axis = 1)
      y = filtered_train_data['Annual_Premium']
```

0.5.1 Train and Test Split

```
[33]: X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.2,
      ↪random_state = 3)
```

```
[34]: print('xtrain', X_train.shape)
      print('ytrain', y_train.shape)
      print('xtest', X_test.shape)
      print('ytest', y_test.shape)
```

```
xtrain (243909, 11)
ytrain (243909,)
xtest (60978, 11)
ytest (60978,)
```

```
[35]: rf1 = RandomForestRegressor(random_state=1)
      rf1.fit(X_train, y_train)
      y_pred1 = rf1.predict(X_test)
      print(np.sqrt(mean_squared_error(y_test, y_pred1)))
```

```
0.04100899619245232
```

```
[36]: lm1 = LinearRegression()
      lm1.fit(X_train, y_train) #lm1.fit(X_train, y_train)
      y_pred = lm1.predict(X_test)
```

```
print(np.sqrt(mean_squared_error(y_test, y_pred)))
```

0.04289074001965307

```
[37]: model=xgb.XGBRegressor(objective='reg:squarederror',learning_rate= 0.1,
    ↪max_depth=7, n_estimators=300)#objective='reg:squarederror',learning_rate= 0.
    ↪1, max_depth=7, n_estimators=300)
model.fit(X_train,y_train)
y_predict=model.predict(X_test)
# print('mean absolute', mean_absolute_error(y_test, y_predict))
mse = mean_squared_error(y_test, y_predict)
# print('mean squared error', mse)
print('root mean squared error', np.sqrt(mse))

print(y_predict)
```

root mean squared error 0.037306059685084846

[0.10049621 0.06478452 0.0703342 ... 0.08173221 0.08770123 0.06526171]

0.5.2 Building ML Model

0.5.3 On full data for competition

```
[38]: lm2 = LinearRegression()
lm2.fit(X, y)
```

```
[38]: LinearRegression()
```

```
[39]: model2=xgb.XGBRegressor(objective='reg:squarederror',learning_rate= 0.1,
    ↪max_depth=7, n_estimators=300)#objective='reg:squarederror',learning_rate= 0.
    ↪1, max_depth=7, n_estimators=300)
model2.fit(X,y)
```

```
[39]: XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
    colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
    importance_type='gain', interaction_constraints='',
    learning_rate=0.1, max_delta_step=0, max_depth=7,
    min_child_weight=1, missing=nan, monotone_constraints='()',
    n_estimators=300, n_jobs=4, num_parallel_tree=1, random_state=0,
    reg_alpha=0, reg_lambda=1, scale_pos_weight=1, subsample=1,
    tree_method='exact', validate_parameters=1, verbosity=None)
```

0.6 Prediction for Test Dataset

```
[40]: #drop the "id" column
test_id = filtered_test_data['id'].to_numpy()
print(test_id)
filtered_test_data = filtered_test_data.drop(['id'], axis = 1)
filtered_test_data
```

```
[ 16222 342481 20916 ... 303222 333830 245084]
```

```
[40]:
```

	Gender	Age	Driving_License	Region_Code	Previously_Insured	\
0	0	0.980501	1	28.0	0	
1	1	-0.826319	1	3.0	0	
2	0	-0.890848	1	28.0	0	
3	0	-0.632731	1	24.0	0	
4	1	-0.697260	1	28.0	0	
...	
76217	0	0.077091	1	28.0	0	
76218	1	0.012562	1	48.0	0	
76219	0	1.367677	1	28.0	0	
76220	0	-1.019907	1	15.0	1	
76221	0	1.238618	1	8.0	1	

	Vehicle_Damage	Policy_Sales_Channel	Vintage	Vehicle_Age_1_2_Year	\
0	1	26.0	-1.209060	0	
1	1	156.0	1.503816	1	
2	1	124.0	1.205041	0	
3	1	157.0	0.966021	1	
4	1	26.0	1.061629	0	
...	
76217	1	124.0	-0.169323	1	
76218	0	124.0	0.284815	1	
76219	1	26.0	1.563571	1	
76220	0	152.0	-0.671265	0	
76221	0	26.0	0.966021	1	

	Vehicle_Age_lt_1_Year	Vehicle_Age_gt_2_Years
0	0	1
1	0	0
2	1	0
3	0	0
4	1	0
...
76217	0	0
76218	0	0
76219	0	0
76220	1	0
76221	0	0

```
[76222 rows x 11 columns]
```

```
[41]: predictions = model2.predict(filtered_test_data)
print(predictions)
```

```
[0.09831095 0.02479893 0.09748651 ... 0.09714258 0.0757018 0.08505565]
```

```
[42]: test_pred = lm2.predict(filtered_test_data)
```

0.7 Saving Prediction File

```
[43]: data = {"id":test_id, "Annual_Premium":predictions}
res= pd.DataFrame(data)
res.index = filtered_test_data.index # for comparision
res.to_csv("prediction_result5.csv", index = False)
```